

## Download the Dataset

<https://www.kaggle.com/code/kredy10/simple-lstm-for-text-classification/data>

## Import required library

```
from keras_preprocessing.sequence import pad_sequences
import pandas as pd
import numpy as np
import re
import collections
import contractions
import seaborn as sns
import matplotlib.pyplot as plt
plt.style.use('dark_background')
from statistics import mode
import nltk
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
import warnings
warnings.simplefilter(action='ignore', category=Warning)
import keras
from keras.layers import Dense, Embedding, LSTM, Dropout
from keras.models import Sequential, load_model
from keras.preprocessing.text import Tokenizer
```

## Read dataset and do pre-processing

```
df = pd.read_csv("spam.csv", encoding='latin-1')
df.head()
```

```
      v1                                                    v2 Unnamed: 2
\
0  ham  Go until jurong point, crazy.. Available only ...      NaN
1  ham                                Ok lar... Joking wif u oni...      NaN
2  spam  Free entry in 2 a wkly comp to win FA Cup fina...      NaN
3  ham  U dun say so early hor... U c already then say...      NaN
4  ham  Nah I don't think he goes to usf, he lives aro...      NaN

      Unnamed: 3  Unnamed: 4
0           NaN           NaN
1           NaN           NaN
```

```

2      NaN      NaN
3      NaN      NaN
4      NaN      NaN

```

```
df.shape
```

```
(5572, 5)
```

```
df.drop(["Unnamed: 2", "Unnamed: 3", "Unnamed: 4"], axis=1,
        inplace=True)
```

```
df.columns = ["Category", "Tweet"]
df.head()
```

```

   Category      Tweet
0      ham  Go until jurong point, crazy.. Available only ...
1      ham                Ok lar... Joking wif u oni...
2     spam  Free entry in 2 a wkly comp to win FA Cup fina...
3      ham  U dun say so early hor... U c already then say...
4      ham  Nah I don't think he goes to usf, he lives aro...

```

```
df.isnull().sum()
```

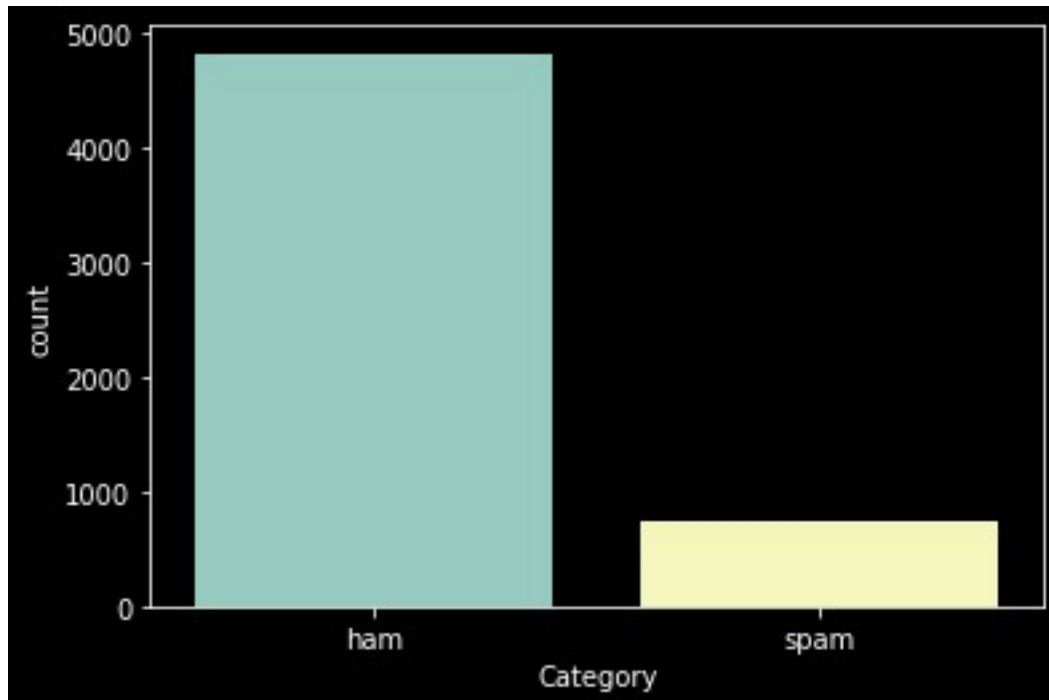
```

Category    0
Tweet       0
dtype: int64

```

```
sns.countplot(x=df['Category'])
```

```
<AxesSubplot:xlabel='Category', ylabel='count'>
```



```

from sklearn.preprocessing import LabelEncoder
en = LabelEncoder()
y = en.fit_transform(df["Category"])

tokenizer = Tokenizer()
tokenizer.fit_on_texts(df['Tweet'])
text_to_sequence = tokenizer.texts_to_sequences(df['Tweet'])

max_length_sequence = max([len(i) for i in text_to_sequence])
padded_sms_sequence = pad_sequences(text_to_sequence,
maxlen=max_length_sequence,
padding = "pre")
TOT_SIZE = len(tokenizer.word_index)+1

```

## Create Model

```
lstm_model = Sequential()
```

## Add Layers (LSTM, Dense-(Hidden Layers), Output)

```

lstm_model.add(Embedding(TOT_SIZE, 32,
input_length=max_length_sequence))
lstm_model.add(LSTM(100))
lstm_model.add(Dropout(0.4))
lstm_model.add(Dense(20, activation="relu"))
lstm_model.add(Dropout(0.3))
lstm_model.add(Dense(1, activation = "sigmoid"))

```

```
lstm_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 189, 32)	285472
lstm (LSTM)	(None, 100)	53200
dropout (Dropout)	(None, 100)	0
dense (Dense)	(None, 20)	2020
dropout_1 (Dropout)	(None, 20)	0
dense_1 (Dense)	(None, 1)	21

Total params: 340,713

Trainable params: 340,713

Non-trainable params: 0

---

## Compile the Model

```
lstm_model.compile(loss = "binary_crossentropy", optimizer = "adam",  
metrics = ["accuracy"])
```

## Fit the Model

```
lstm_model.fit(padded_sms_sequence, y, epochs = 10,  
validation_split=0.2, batch_size=16)
```

```
Epoch 1/10  
279/279 [=====] - 40s 133ms/step - loss:  
0.1688 - accuracy: 0.9522 - val_loss: 0.0606 - val_accuracy: 0.9857  
Epoch 2/10  
279/279 [=====] - 37s 131ms/step - loss:  
0.0326 - accuracy: 0.9917 - val_loss: 0.0590 - val_accuracy: 0.9865  
Epoch 3/10  
279/279 [=====] - 37s 132ms/step - loss:  
0.0108 - accuracy: 0.9975 - val_loss: 0.0480 - val_accuracy: 0.9874  
Epoch 4/10  
279/279 [=====] - 37s 132ms/step - loss:  
0.0068 - accuracy: 0.9991 - val_loss: 0.0970 - val_accuracy: 0.9776  
Epoch 5/10  
279/279 [=====] - 37s 131ms/step - loss:  
0.0076 - accuracy: 0.9978 - val_loss: 0.0762 - val_accuracy: 0.9803  
Epoch 6/10  
279/279 [=====] - 37s 132ms/step - loss:  
0.0055 - accuracy: 0.9993 - val_loss: 0.0489 - val_accuracy: 0.9874  
Epoch 7/10  
279/279 [=====] - 37s 131ms/step - loss:  
0.0018 - accuracy: 1.0000 - val_loss: 0.0634 - val_accuracy: 0.9848  
Epoch 8/10  
279/279 [=====] - 37s 132ms/step - loss:  
2.8042e-04 - accuracy: 1.0000 - val_loss: 0.0760 - val_accuracy:  
0.9901  
Epoch 9/10  
279/279 [=====] - 39s 140ms/step - loss:  
3.0002e-04 - accuracy: 1.0000 - val_loss: 0.0797 - val_accuracy:  
0.9883  
Epoch 10/10  
279/279 [=====] - 39s 138ms/step - loss:  
1.4242e-04 - accuracy: 1.0000 - val_loss: 0.0906 - val_accuracy:  
0.9892
```

<keras.callbacks.History at 0x1beadc4a380>

## Save The Model

```
lstm_model.save(r'D:\IBM Project\spam_classifier.h5')
```

## Test The Model

```
model=load_model(r'D:\IBM Project\spam_classifier.h5')

text='Urgent UR awarded a complimentary trip to EuroDisinc Trav,
Aco&Entry41 Or å£1000. To claim txt DIS to 87121
18+6*å£1.50(moreFrmMob. ShrAcomOrSglSuplt)10, LS1 3AJ'
text = tokenizer.texts_to_sequences(text)
max_length = max([len(i) for i in text])
padded_sms = pad_sequences(text, maxlen=max_length,
                           padding = "pre")

prediction = np.squeeze(model.predict(padded_sms))
classes_x = (prediction > 0.5).astype("int32")

mode_val=mode(classes_x)
if(mode_val==1):
    print("Ham")
else:
    print("Spam")

6/6 [=====] - 0s 4ms/step
Spam
```